

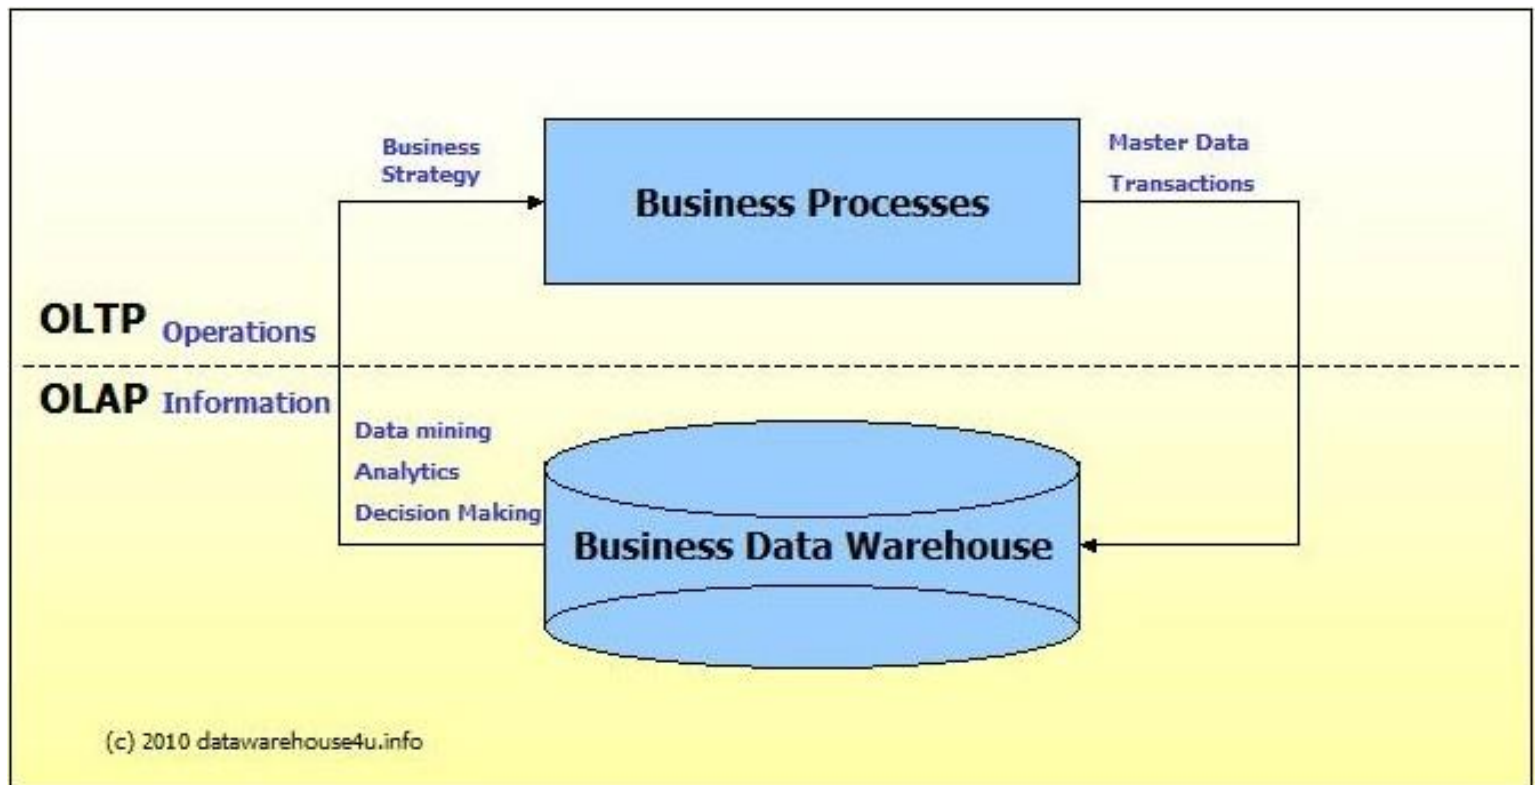


NoSQL


W2013  
CSCI 2141

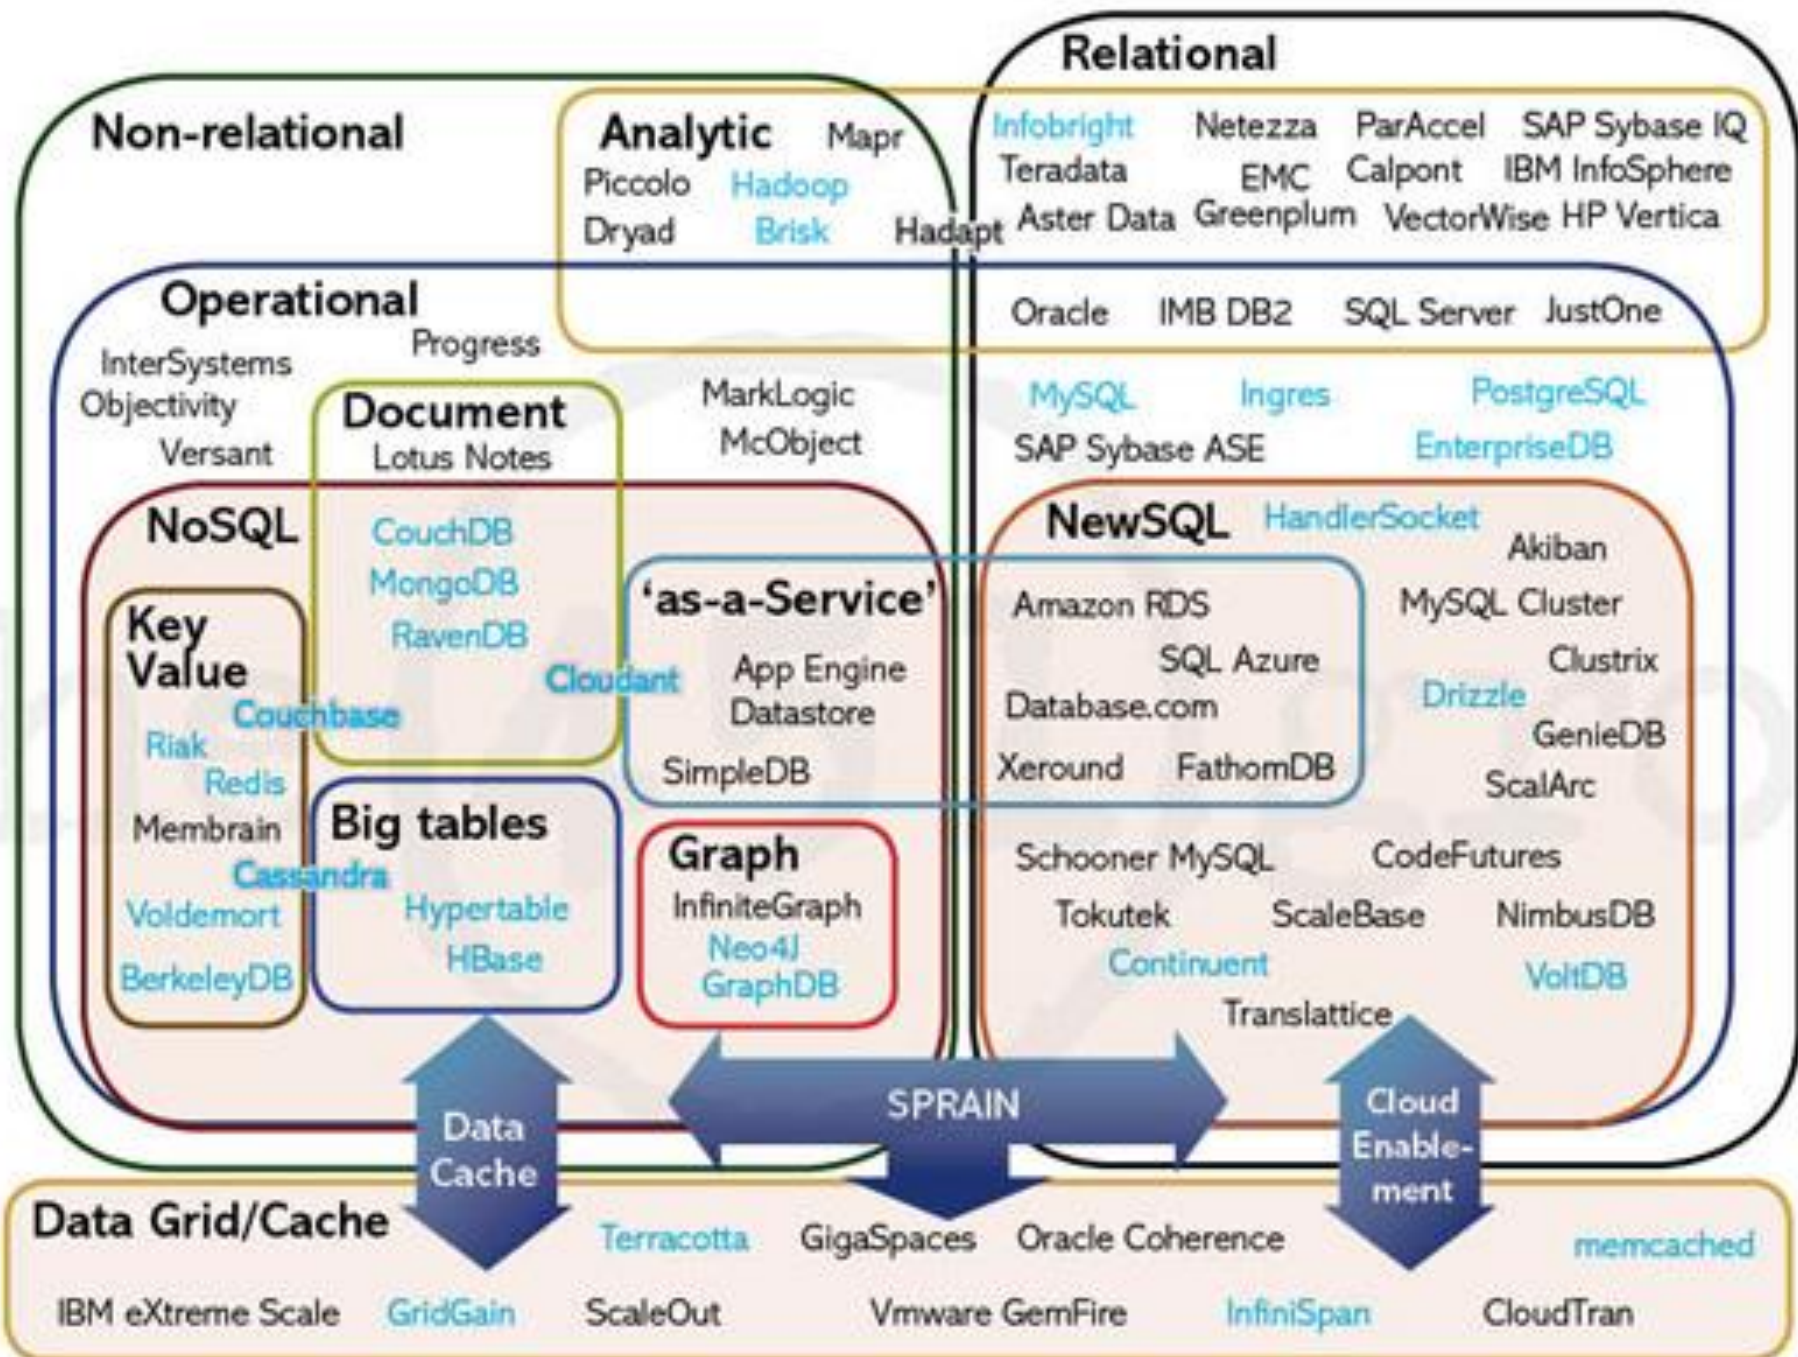
# + OLTP vs. OLAP

- We can divide IT systems into transactional (OLTP) and analytical (OLAP). In general we can assume that OLTP systems provide source data to data warehouses, whereas OLAP systems help to analyze it



# + Challenges of Scale Differ







# A Comparison of SQL and NoSQL Databases

Slides from: Keith W. Hare  
Metadata Open Forum

More reading: <http://martinfowler.com/articles/nosqlKeyPoints.html>

Metadata Open Forum




# Abstract

NoSQL databases (either no-SQL or Not Only SQL) are currently a hot topic in some parts of computing. In fact, one website lists over a hundred different NoSQL databases.

This presentation reviews the features common to the NoSQL databases and compares those features to the features and capabilities of SQL databases.

# BIG DATA!

# HOW TO WRITE A CV



Leverage the NoSQL boom

21 March 2016

# + SQL Characteristics

- Data stored in columns and tables
- Relationships represented by data
- Data Manipulation Language
- Data Definition Language
- Transactions
- Abstraction from physical layer



# + SQL Physical Layer Abstraction

- **Applications specify what, not how**
- **Query optimization engine**
- **Physical layer can change without modifying applications**
  - **Create indexes to support queries**
  - **In Memory databases**



# Data Manipulation Language (DML)



- Data manipulated with Select, Insert, Update, & Delete statements
  - Select T1.Column1, T2.Column2 ...  
From Table1, Table2 ...  
Where T1.Column1 = T2.Column1 ...
- Data Aggregation
- Compound statements
- Functions and Procedures
- Explicit transaction control

# + Data Definition Language

- Schema defined at the start
- Create Table (Column1 Datatype1, Column2 Datatype 2, ...)
- Constraints to define and enforce relationships
  - Primary Key
  - Foreign Key
  - Etc.
- Triggers to respond to Insert, Update , & Delete
- Stored Modules
- Alter ...
- Drop ...
- Security and Access Control



# Transactions – ACID Properties



- Atomic – All of the work in a transaction completes (commit) or none of it completes
- Consistent – A transaction transforms the database from one consistent state to another consistent state. Consistency is defined in terms of constraints.
- Isolated – The results of any changes made during a transaction are not visible until the transaction has committed.
- Durable – The results of a committed transaction survive failures



# NewSQL: more OLTP throughput, real-time analytics

- 1) SQL as the primary mechanism for application interaction
- 2) ACID support for transactions
- 3) A non-locking concurrency control mechanism so real-time reads will not conflict with writes, and thereby cause them to stall.
- 4) An architecture providing much higher per-node performance than available from the traditional "elephants"
- 5) A scale-out, shared-nothing architecture, capable of running on a large number of nodes without bottlenecking

# + NoSQL Definition

From [www.nosql-database.org](http://www.nosql-database.org):

Next Generation Databases mostly addressing some of the points: being non-relational, distributed, open-source and horizontal scalable. The original intention has been modern web-scale databases. The movement began early 2009 and is growing rapidly. Often more characteristics apply as: schema-free, easy replication support, simple API, eventually consistent / BASE (not ACID), a huge data amount, and more.

# + NoSQL Products/Projects


<http://www.nosql-database.org/>  
lists 122 NoSQL Databases

- Cassandra
- CouchDB
- Hadoop & Hbase
- MongoDB
- StupidDB
- Etc.

# + NoSQL Produ

<http://www.nosql-d>  
lists 122 NoSQL Dat

- Cassandra
- CouchDB
- Hadoop & Hbase
- MongoDB
- StupidDB
- Etc.





# + NoSQL Distinguishing Characteristics

- Large data volumes
  - Google's "big data"
- Scalable replication and distribution
  - Potentially thousands of machines
  - Potentially distributed around the world
- Queries need to return answers quickly
- Mostly query, few updates
- Asynchronous Inserts & Updates
- Schema-less
- ACID transaction properties are not needed – BASE
- CAP Theorem
- Open source development

# + BASE Transactions

## ■ Acronym contrived to be the opposite of ACID

- Basically Available,
- Soft state,
- Eventually Consistent

## ■ Characteristics

- Weak consistency – stale data OK
- Availability first
- Best effort
- Approximate answers OK
- Aggressive (optimistic)
- Simpler and faster

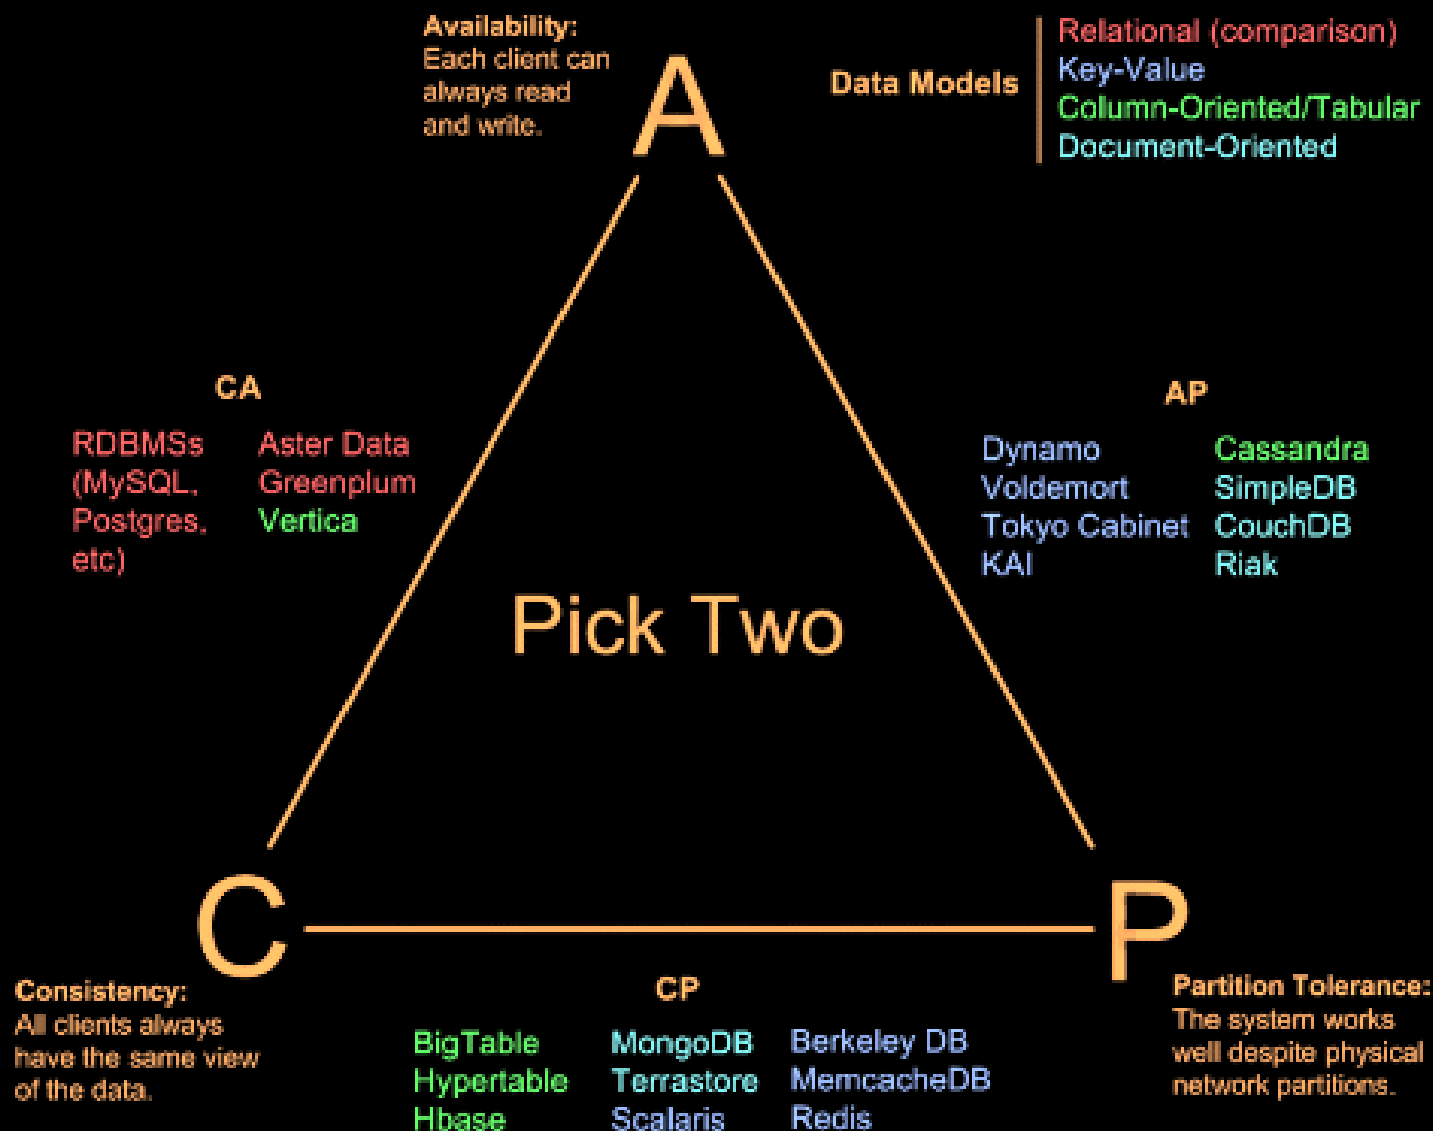
# + Brewer's CAP Theorem

A distributed system can support only two of the following characteristics:

- Consistency
- Availability
- Partition tolerance



# Visual Guide to NoSQL Systems



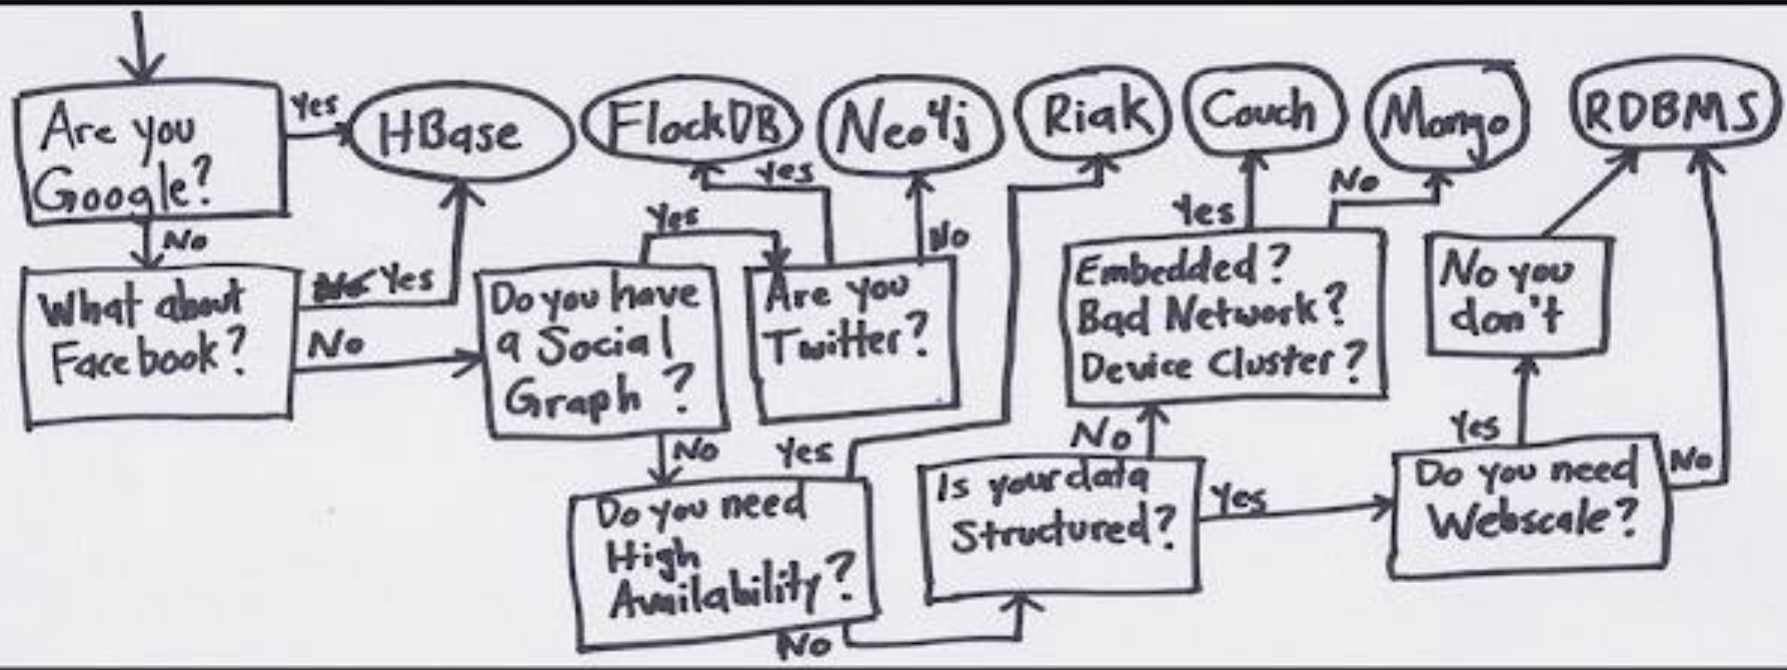
# + NoSQL Database Types

Discussing NoSQL databases is complicated because there are a variety of types:

- Column Store – Each storage block contains data from only one column
- Document Store – stores documents made up of tagged elements
- Key-Value Store – Hash table of keys

# + Other Non-SQL Databases

- XML Databases
- Graph Databases
- Coddasy1 Databases
- Object Oriented Databases
- Etc...
- Will not address these today





Key-Value



Key Value



Ordered Key-Value



Big Table




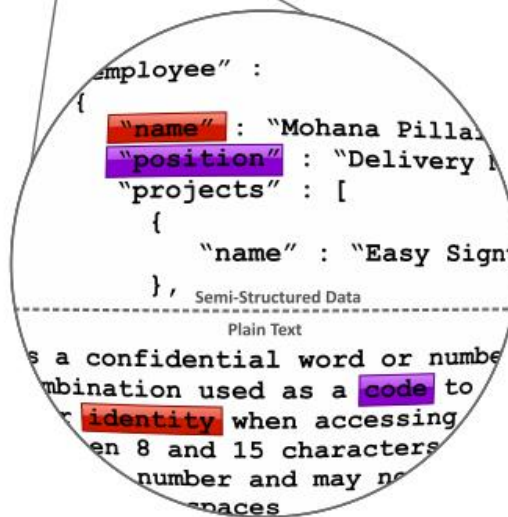
Document,  
Full-Text Search



Graph



SQL





# + Storing and Modifying Data

- Syntax varies
  - HTML
  - Java Script
  - Etc.
- Asynchronous – Inserts and updates do not wait for confirmation
- Versioned
- Optimistic Concurrency

# + Retrieving Data

## ■ Syntax Varies

- No set-based query language
- Procedural program languages such as Java, C, etc.

## ■ Application specifies retrieval path

## ■ No query optimizer

## ■ Quick answer is important

## ■ May not be a single “right” answer

# + Open Source

- Small upfront software costs
- Suitable for large scale distribution on commodity hardware

# + NoSQL Summary

- NoSQL databases reject:
  - Overhead of ACID transactions
  - “Complexity” of SQL
  - Burden of up-front schema design
  - Declarative query expression
  - Yesterday’s technology
- Programmer responsible for
  - Step-by-step procedural language
  - Navigating access path

# + Summary

## ■ SQL Databases

- Predefined Schema
- Standard definition and interface language
- Tight consistency
- Well defined semantics

## ■ NoSQL Database

- No predefined Schema
- Per-product definition and interface language
- Getting an answer quickly is more important than getting a correct answer

# + Web References

- “NoSQL -- Your Ultimate Guide to the Non - Relational Universe!”  
<http://nosql-database.org/links.html>
- “NoSQL (RDBMS)”  
<http://en.wikipedia.org/wiki/NoSQL>
- PODC Keynote, July 19, 2000. *Towards Robust. Distributed Systems*. Dr. Eric A. Brewer. Professor, UC Berkeley. Co-Founder & Chief Scientist, Inktomi.  
[www.eecs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf](http://www.eecs.berkeley.edu/~brewer/cs262b-2004/PODC-keynote.pdf)
- “Brewer's CAP Theorem” posted by Julian Browne, January 11, 2009.  
<http://www.julianbrowne.com/article/viewer/brewers-cap-theorem>
- “How to write a CV” Geek & Poke Cartoon  
<http://geekandpoke.typepad.com/geekandpoke/2011/01/nosql.html>



# Web References

- “Exploring CouchDB: A document-oriented database for Web applications”, Joe Lennon, Software developer, Core International.  
<http://www.ibm.com/developerworks/opensource/library/os-couchdb/index.html>
- “Graph Databases, NOSQL and Neo4j” Posted by Peter Neubauer on May 12, 2010 at:  
<http://www.infoq.com/articles/graph-nosql-neo4j>
- “Cassandra vs MongoDB vs CouchDB vs Redis vs Riak vs HBase comparison”, Kristóf Kovács.  
<http://kkovacs.eu/cassandra-vs-mongodb-vs-couchdb-vs-redis>
- “Distinguishing Two Major Types of Column-Stores” Posted by Daniel Abadi on March 29, 2010  
[http://dbmsmusings.blogspot.com/2010/03/distinguishing-two-major-types-of\\_29.html](http://dbmsmusings.blogspot.com/2010/03/distinguishing-two-major-types-of_29.html)

# + Web References

- “MapReduce: Simplified Data Processing on Large Clusters”, Jeffrey Dean and Sanjay Ghemawat, December 2004.  
<http://labs.google.com/papers/mapreduce.html>
- “Scalable SQL”, ACM Queue, Michael Rys, April 19, 2011  
<http://queue.acm.org/detail.cfm?id=1971597>
- “a practical guide to noSQL”, Posted by Denise Miura on March 17, 2011 at  
<http://blogs.marklogic.com/2011/03/17/a-practical-guide-to-nosql/>